

# COMP2221 Networks

David Head

University of Leeds

Lecture 9

## Reminder of the Last Lectures


In the last two lectures we have seen how to implement a client and a simple server in Java.

The **client** uses the `Socket` class to:

- Contact the server on a prescribed host **and port**.
- Communicate according to an agreed **protocol**.

The **server** uses the `ServerSocket` class to:

- **Listen** to the prescribed port by calling `accept()`.
- Once a connection is made, `accept()` returns a `Socket` object that the server can use to implement its part of the protocol.



# Problems with the simple server

We also mentioned some **limitations** of the simpler server:

`accept()` is **blocking**:

- It only returns once a client makes contact.
- In the meantime, the **entire server application is doing nothing**.

Only communicates with **one client at a time**:

- Other clients will get **queued**, leading to delays.
- Server may be frequently **idle** during communication with the client (*i.e.* while waiting for a response).
- Only one **protocol** handler per server, rather than one per client (e.g. `KnockKnockProtocol`).

# Today's lecture

Whereas creating one protocol handler per client is conceptually simple, the other problems are more serious.

There are two solutions:

- ① Use **non-blocking communication** as provided by `java.nio`.
- ② Use **concurrent threads**, with one or more clients per thread.

We will consider (2) over the next 3 lectures, and (1) in lecture 12.

# Why parallel computation?

Today is the first of 3 lectures when we will consider **parallel computation**.

- *i.e.* writing software that makes use of hardware capable of performing calculations 'in parallel,' *i.e.* **simultaneously**.
- An increasingly important skill, as almost all modern hardware **is** parallel.


For the BCS-accredited programmes, you are required to cover some parallel programming.

- For some of you, this is the **only** time you will encounter parallel programming in your programme.
- Some optional Level 3 modules (*esp.* COMP3221) cover this much more extensively.

# The rise of parallel architectures


In 1965 Gordon Moore made the empirical observation that the density of transistors on a chip doubles every 18-24 months.

- Plotted component cost vs. no. of components.
- Included projected 1970 data.
- Identified **exponential increase** of the most cost-effective number of components.



Although the law still holds, processor speeds no longer track it:

### Stuttering




<http://www.economist.com/technology-quarterly/2016-03-12/after-moores-law>

The problem is heat generation increases rapidly with clock speed.



# The benefits of parallelism

One way to improve performance is to allow calculations to be performed **simultaneously**, *i.e.* in **parallel**:



# The limits of automated parallelism

Chip designers have already tried to automate some parallelism.

**Instruction-level parallelism** (ILP) processors use a *pipelining* architecture in which each instruction is processed in multiple stages, somewhat like an assembly line at a factory.

- Different stages for subsequent instructions can be performed simultaneously.

CPUs also have **multiple functional units** such as FPUs, ALUs *etc.* that can work independently.

- **Superscalar** architectures can merge nearby calculations into a single one.

However, each of these have limits - they do not **scale**.

# Current parallel architectures

Almost all modern architectures contain multiple **processing units**, e.g. multiple cores, CPUs *etc.*

**Pro:** **Scales** better than previous automated parallelism.

**Con:** Software must be **specifically developed** to take into account these processing units.

Current parallel hardware can be broadly classified into:


- ① Shared memory architectures.
- ② Distributed memory architectures.
- ③ Graphics Processing Units or **GPUs**.

# 1. Shared memory architectures

## Definition

All **processing units** (e.g. *cores*) access the same memory.

This includes anything with one or more **multi-core CPUs**, so almost all modern desktops, laptops, tablets, smart phones *etc.*




## 2. Distributed memory architectures

### Definition

Processing units only access a **fraction** of total memory available


Includes **High-Performance Computing** (HPC) clusters (e.g. supercomputers), and **distributed systems** ('cloud computing').



**Benchmark  
speeds of all  
supercomputers in  
the world [Nov.  
2024].**

1 EFlop/s  
= 1 **exa-FLOP**  
=  $10^{18}$  floating  
point operations  
per second.

*[From top500.org]*



### 3. Graphics Processing Units (GPUs)

#### GPUs

Hardware specifically designed to rapidly perform calculations that arise in the graphical rendering of scenes

Historically driven by graphics applications, especially video games, but increasing being used for **non-graphical** applications such as **machine learning, cryptocurrencies** *etc.*

- Some devices design for general-purpose calculations.
- Known as **GPGPUs** for General-Purpose GPUs.

GPU hardware has multiple memory stores, some of which can be viewed as distributed, some shared.

# The most powerful 'computer' known?

Arguably the most complex system known is the **human brain**.

If regarded as a computer, it would be **massively parallel**:

- Synapse speeds are about 5 ms, so the 'clock speed' would be less than 1kHz.
- We have about  $10^{11}$  neurons, each connected to  $10^4$  others.
- The current fastest supercomputer has  $10^7$  cores.



<http://scitechconnect.elsevier.com>



# How to program parallel architectures?

Parallel architectures allow multiple calculations to be performed **simultaneously**.

- Can improve performance without increasing the clock speed.

However, what if one calculation requires the result of *another* calculation as input? How can they be performed *in parallel*?<sup>1</sup>

Algorithms typically need to be **redesigned** to make good use of parallel hardware. These new algorithms need to be designed and implemented by **programmers, i.e. us!**

---

<sup>1</sup>Answer: They can't ...

# Concurrency

**Concurrency** is when two or more tasks are in progress in the same **time frame**<sup>1</sup>.

For instance, for an event-driven GUI, a user event (e.g. a mouse click) might result in a **callback function** being called.

- Would normally be called by a separate task (*thread*), so that the main loop can continue (and detect other user events).

Can be achieved using **interrupts**:

- OS 'slices' CPU time amongst all running tasks.
- Most commonly **pre-emptive** multi-tasking.

---

<sup>1</sup>McCool et. al., *Structured Parallel Programming* (Morgan-Kaufman, 2012).

# Parallelism

An increasingly important concept related to concurrency is **parallelism**.

**Parallelism** refers to the ability to perform **multiple calculations simultaneously** by using more than one **processing unit**.

Possible to be concurrent but *not* parallel:

- Multi-tasking is possible on a single-core CPU, which only has one compute unit (*i.e.* the core).

Parallel  $\implies$  Concurrent **but** Concurrent  $\nRightarrow$  Parallel

# Processes *versus* threads

There are two basic units of execution, **processes** and **threads**.

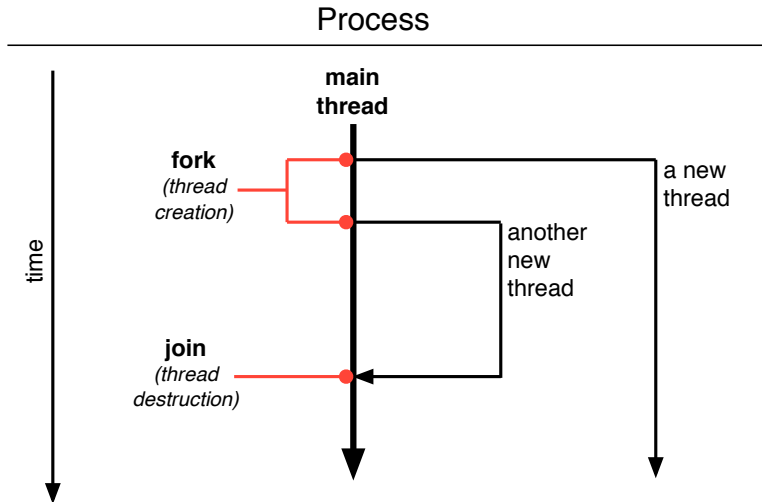
## Process:

- A **self-contained** execution environment.
- Private set of run-time resources.
- Has its own **heap** memory.
  - All objects created with `new` (inc. arrays), global variables *etc.*
- Also has its own **stack** memory, which has local variables, function arguments *etc.*
- **Expensive** to create and destroy.
- Normally 1 application = 1 process, but not always:
  - For example, each **tab** in an application may correspond to a different process

## Threads:

- Are **launched by**, and **exist within**, a process.
- Every process has at least one (the **main** thread in Java).
- Have their own stack memory.
- **Shares** the heap memory with the launching process.
  - Potential problems if multiple concurrent threads can read and write the same data.
- Is **lightweight**; small cost for creating and destroying threads.
- How threads are assigned to cores is up to the OS, more specifically the **scheduler**.

# One process, multiple threads



# Parallel languages and frameworks

Most languages support parallel programming at some level:

- C, C++ has libraries/standards, e.g. OpenMP, Cilk, the pthread library, etc.
- C++11 has native multi-threading support.
- Python has a threading library<sup>1</sup>.

Even early versions of Java supported multi-threading at the **language level**.

- At least partly motivated by event-driven GUI design.
- *i.e.* `java.lang.Thread`, `java.util.concurrent`.

---

<sup>1</sup>Although need to work around its Global Interpreter Lock (GIL) to exploit multi-cores.

# Summary and next time

Today we have looked at parallel architectures and software.

- Why parallel computation is becoming increasingly important.
- Shared *versus* distributed memory.
- Concurrency *versus* parallelism.
- Processes *versus* threads.

Over the next two lectures we will various ways of implementing a multi-threaded server in Java.