

COMP2221 Networks

David Head

University of Leeds

Lecture 14

Previous lectures

For all of our client-server examples so far we have used `Socket`, `ServerSocket`, or classes derived from them (*i.e.* `SSLSocket`).

- Sits in the Application layer.
- Use the Transport layer protocol TCP = Transmission Control Protocol.
- Creates a **persistent** contact between the two hosts.
- The most common **protocol** for sending and receiving messages over the internet.

Today's lecture

Today we will look at 'the other' protocol: UDP = User Datagram Protocol.

- A **connectionless** protocol, *i.e.* there is no **persistent** contact.
- Handled in Java using two new classes in `java.net`:
 - `DatagramPacket`
 - `DatagramSocket`
- See an example of a UDP-based client-server application.

We will look in more detail at the Transport-layer details of UDP (and TCP) in Lecture 16.

- In this lecture we focus on the protocol, *i.e.* what the Application layer 'needs to know.'

TCP

Transmission Control Protocol:

- **Reliable:** Lost/damaged packets are automatically re-sent.
- Re-orders data to maintain **sequencing**.
- **Throttles** the connection to avoid packet loss.
- Slower than UDP.

Analogy: **Phone call**

- Connect two people first.
- Maintain the connection throughout the conversation.
- Ordered communication.

UDP

User Datagram Protocol:

- **Unreliable:** No guarantee of delivery.
- Data is accepted **in the order that it arrives**.
- No congestion control.
- Faster than TCP.

Analogy: **email**

- Messages do not necessarily arrive in the order they were sent.
- No persistent contact between sender and receiver.
- No guarantee the mail will arrive at all.

When to use UDP

UDP is not suitable for ftp/http-type applications which require **complete** and **ordered** data.

- Reliability is a priority here.

UDP is suitable for applications where **speed** is a priority.

- e.g. streaming audio/video, where a small fraction of lost packets is acceptable.
- DNS uses UDP (see Lecture 4).

UDP is also suitable for:

- Testing for reliability, *i.e.* send a UDP packet and see if it is returned within a certain time.
- **Multi-casting** (see next Lecture).

Real-time networking

If not all of the services for TCP are necessary, can start with UDP and add required services **in the Application layer**.

A common application is for **real-time conversation**:

- Need **rapid response** — no more than 400ms.
- Can still follow the conversation with **some** packet loss.
- Need some strategy to **recover** from packet loss — including if it arrives 'too late.'

Common examples include:

- **VoIP** — Voice-over-IP.
- **RTP** — Real-time Transport Protocol.

RTP = Reat-time Transport Protocol

Open standard for **real-time conversational applications**:

- Runs on-top of UDP (typically) — seen as any other UDP packet in the Network layer.
- 12-byte **RTP header** — sequence numbers, time-stamps, *etc.*
- Better chance of interacting if both end applications use RTP.

There are also **proprietary protocols** such as **Skype**:

- Uses UDP for audio/video data packets, TCP for control.
- Media packets also sent over TCP if a firewall blocks UDP (*see Lecture 17*).

For more on media streaming, see Kurose and Ross 7th ed., Chapter 9.

UDP in Java

The implementation is split into two parts, both in `java.net`:

The `DatagramPacket` class:

- Loads and unloads data into a **datagram**.
- The destination **address** and **port** is part of the datagram — *not* the socket!
- The source address and port are added automatically.

The `DatagramSocket` class:

- Sends or receives a datagram.
- Only knows the **local port** on which it listens or sends.

Issues for the client-server model

TCP treats a network connection as a **stream**.

- Permanent, two-way connection.
- We can assume sent data is received.
- We can assume it arrives in the order it was sent.

UDP has no concept of a unique, permanent connection between two hosts.

- Only deals with single messages/packets.
- Each host has to **listen** for data.
- The DatagramSocket does not know the destination.
- Does **not** require a one-to-one connection; can be e.g. one-to-many [*cf. Lecture 15*].

The DatagramPacket class

Two types of constructor depending on context of use:

For **receiving** a datagram:

```
public DatagramPacket( byte[] buffer, int length )
```

- Maximum length specified by protocol; 8K is typical.
- Buffer capacity `buffer.length` must be **at least** as large as the `length` argument.

For **sending** a datagram:

```
public DatagramPacket( byte[] data, int length, InetAddress  
destination, int port )
```

- Data is **loaded** from the given array.
- Destination address and port included.

Getters

```
public byte[] getData()
```

- Returns the data buffer as a byte array.
- **Not** an I/O stream — **we** convert to/from bytes.
- Up to us to break larger messages into manageable 'chunks.'

```
public int getLength()
```

- Length of data to be sent, or received.

```
public InetAddress getAddress()
```

- Returns the IP address of the destination (if sending) or the source (if receiving).

```
public int getPort()
```

- The destination or source port.

Setters

Unlike Sockets, we **can** change the class fields after construction.

- Allowed because there is no persistent connection.
- Can be useful if we want to re-use a datagram, to avoid garbage collection (*i.e.* for performance).

```
public void setData( byte[] buffer )  
public void setLength( int length )  
public void setAddress( InetAddress address )  
public void setPort( int port )
```

The DatagramSocket class

Clients and servers use the same class (*i.e.* there is no 'DatagramServerSocket'), but different constructors.

For a **client**:

`public DatagramSocket()` throws `SocketException`

- Opens a port (assigned at run time).

For a **server**:

`public DatagramSocket(int port)` throws `SocketException`

- Opens on a **given** port (which is published for the client).

Both are **bound** to a local port and **listen** for incoming data.

Useful methods (1)

```
public int getLocalPort()
```

- Returns the port to which the socket is bound.
- Only really useful for the client-type constructor.

```
public InetAddress getLocalAddress()
```

- Gets the local address to which the socket is bound.
- Can be useful for hardware with multiple IP addresses.

```
public void close()
```

- Closes the DatagramSocket.

Useful methods (2)

```
public void send( DatagramPacket d ) throws IOException
```

- Sends the datagram you have created.

```
public void receive( DatagramPacket d ) throws IOException
```

- Receives a single datagram and stores it in d.
- **Blocking** - does not return until data is received.
- Closest thing to `accept()` in TCP's `ServerSocket`.
- Can use multi-threaded approach similar to before.

```
public void setSoTimeout( int timeout ) throws SocketException
```

- Sets the maximum time (in milliseconds) the socket will block for before throwing a `SocketTimeoutException`.

Convenience methods for managing connections

```
public void connect( InetAddress host, int port )
```

- After calling, can only communicate with the specified destination — will throw an `IllegalArgumentException` if an attempt is made to send to a different destination.
- Will not perform security checks (if security enabled).
- This is not a connection in the TCP sense, but does establish who they communicate with — a **one-sided** connect.

```
public void disconnect()
```

- Breaks the connection, defaults back to general communication (*i.e.* anyone can receive again).

Do not **need** to use these — the example given next doesn't.

Example¹: Echo client/server

Code on Minerva: `UDPEchoClient.java`, `UDPEchoServer.java`

A UDP client/server pair.

The port number for the server is set when launched (command line argument).

The client sends a packet containing a single string.

- Destination address, port and string all command line arguments.

Servers responds with a packet containing the same string.

Client handles possible packet loss *and* rogue connections.

¹cs.baylor.edu/~donahoo/practical/JavaSockets/textcode.html, from the book *TCP/IP Sockets in Java*, Calvert and Donahoo (Morgan-Kaufman, 2001).

UDPEchoServer fragment

```
1 // Bind to port
2 DatagramSocket socket = new DatagramSocket(servPort);
3
4 // Re-use the same data packet
5 DatagramPacket packet = new DatagramPacket(new byte[ECHOMAX
    ],ECHOMAX);
6
7 // Server loop
8 while( true ) {
9     socket.receive(packet);          // Receive packet from client
10
11     // Message to stdout on server
12     System.out.println("Handling client at " +
13         packet.getAddress().getHostAddress() +
14         " on port " + packet.getPort());
15
16     // Send message back to client (i.e. echo)
17     socket.send(packet);
18 }
```

UDPEchoClient fragment (1)

```
1 DatagramSocket socket = new DatagramSocket();
2
3 // Maximum receive blocking time (milliseconds)
4 socket.setSoTimeout(TIMEOUT);
5
6 // The packet to send (address and port specified in args)
7 DatagramPacket sendPacket = new DatagramPacket(bytesToSend,
8         bytesToSend.length, serverAddress, servPort);
9
10 // The packet to receive (should match the one sent
11 // for an echo server)
12 DatagramPacket receivePacket =
13     new DatagramPacket(new byte[bytesToSend.length],
14         bytesToSend.length);
```

UDPEchoClient fragment (2)

```
1 do {
2     socket.send(sendPacket);           // Send the string
3     try {
4         socket.receive(receivePacket);
5
6         // Check for rogue packets
7         if (!receivePacket.getAddress().equals(serverAddress))
8             throw new IOException("Received from unknown source");
9
10        receivedResponse = true;
11    } catch (InterruptedException e) { // Timeout; Retry?
12        System.out.println("Timed out...");
13    }
14 } while( !receivedResponse );
15
16 if( receivedResponse )                // Success or failure?
17     System.out.println("Received: " +
18         new String(receivePacket.getData()));
```

Testing on a real network

This example — or any other client-server application — will never suffer packet loss when tested on the same host, *i.e.* localhost.

- Also very unlikely on a Local Area Network (LAN).

However, when testing over a Wide Area Network (WAN), including the internet, both packet loss and packet corruption become a real possibility.

- We will see one way **how** this can happen in Lecture 17.

Don't be fooled into thinking UDP is simpler than TCP!

Once packet loss is factored in, your program logic will typically become *much* more complex.

Overview and next lecture

Today we have looked at UDP:

- Supports fast but unreliable transfer.
- Implemented in Java with `DatagramPacket` and `DatagramSocket`.
- Seen a simple **one-to-one** example, a UDP echo server.

The code is available on Minerva.

Next time we will see how UDP can be used for **one-to-many** communication.