

# COMP2221 Networks

David Head

University of Leeds

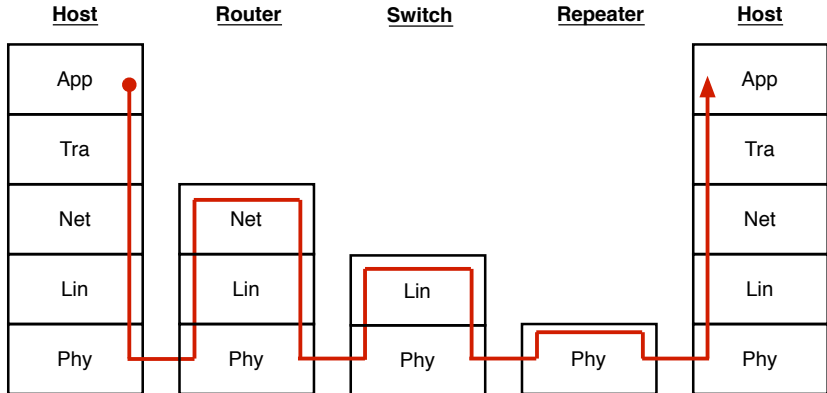
Lecture 20

# Today's lecture

Today's lecture is intended to overview and revise the course in preparation for the exam:

- The exam format.
- What to expect in the exam.
- Summarise course material.
- Leave a few minutes at the end for module feedback.

# The TCP/IP protocol stack (Lecture 2)



After a general introduction, progressed downwards through the 5-layer TCP/IP **protocol stack**:

- **Application** layer for **single processes** (Lectures 5-15).
- **Transport** layer that provides **ports** to control data transfer between processes (Lectures 3, 16).
- **Network** layer that handles **IP Addresses** for hosts and routers (Lectures 4, 5, 17, 18).
- **Link** layer responsible for data communication between adjacent nodes (Lecture 19).
- **Physical** layer responsible for sending individual bits (mentioned in Lecture 19).

Also described the **7-layer OSI stack** (Lecture 2).

# Terminology and addresses

Packets of data are typically given different names for each layer.

All except Physical and Application layers also have a form of 'address'. Source and destination addresses appear in their headers.

Layer	Usual name for packet	Address or similar
Application	Message	-
Transport	Segment <sup>1</sup>	Port
Network	Datagram <sup>1</sup>	IP <sup>2</sup>
Link	Frame	MAC
Physical	-	-

<sup>1</sup>UDP = User **D**atagram Protocol exists in the Transport layer.

<sup>2</sup>Converted to **hostnames** using the Domain Name System.

## Need for buffering

Three layers (Transport, Network and Link) **add headers** (and possibly trailers/footers) to messages.

- Sizes depending on the protocol (TCP/UDP; IPv4/IPv6; any number of Ethernet/IEEE protocols), but can be 20-40 bytes per header level.

This means that the message finally sent over the Physical layer can be *much* larger than the Application data.

- e.g. `telnet` only sends single characters.

This highlights the importance of **buffering** at the Application level to improve I/O performance.

- Conveniently implemented in Java using **streams** (Lecture 6).

# Performance

Other ways to improve performance, in particular for servers:

1. Use **non-blocking I/O** (Lecture 12).
  - Allows a single server thread to deal with multiple clients.
  - Not easy to implement.
2. Use **multi-threading** (Lectures 9-11).
  - Each client handler has its own thread.
  - Easy to implement, especially with Executor's thread pools.
  - Makes good use of modern, multi-core architectures.
  - Strategy applicable to most problems, not just I/O.

# UDP

We also looked at the UDP (Lecture 14).

- **Connectionless**, unlike TCP where a connection is maintained until closed.
- **Unreliable**, so packets may be lost or arrive in a different order than sent (again, unlike TCP).

UDP could be considered for **streaming applications**.

UDP can also be used for **multicasting** (Lecture 15):

- Same data sent to multiple receivers without causing congestion.
- Not widely used, although supported by IPv6.



# Security

We touched on **security**, an increasing important issue.

- Looked at SSLSockets in Lecture 13.
- The basics of **encryption** and **authentication**.
- **Symmetric** key algorithms, which use a **private key**.
- **Asymmetric** key algorithms, which use a private and **public key**.

This important topic is covered fully in the Level 3 module  
*COMP3911 Secure Computing*.

# Lecture summary (1)

- Lecture 1 : Introduction to networks; also admin.
- Lecture 2 : Network architectures; 5 and 7-layer models.
- Lecture 3 : Ports, UDP and TCP; headers for lower levels.
- Lecture 4 : DNS and how it maps names to IP addresses.
- Lecture 5 : IP addresses, IPv4 and IPv6; CIDR and NAT; `InetAddress` in Java.
- Lecture 6 : Java I/O streams, including buffering and filters.
- Lecture 7 : The `Socket` class; construction involves internet access and binding to a port.
- Lecture 8 : The `ServerSocket` class, and how instances listen to ports.

## Lecture summary (2)

- Lecture 9 : Parallel and concurrent programming in general.
- Lecture 10 : The Java Thread class; synchronisation; thread-per-client servers.
- Lecture 11 : Thread pool servers; Executor service.
- Lecture 12 : Non-blocking I/O; Buffer, Channel and Selector.
- Lecture 13 : Network security: Encryption and authentication; the SSLSocket class.
- Lecture 14 : UDP: DatagramPacket and DatagramSocket.
- Lecture 15 : One-to-many communication; multicasting with UDP; the MultiSocket class.

## Lecture summary (3)

- Lecture 16 : Transport layer: Connection management and congestion control; finite state machines; UDP and TCP headers.
- Lecture 17 : Network layer: More details on CIDR; tunnelling; routers, switching fabrics and (generalised) forwarding tables.
- Lecture 18 : Routing algorithms: Dijkstra's algorithm, RIP and advertisements, hierarchical OSPF and BGP.
- Lecture 19 : Link layer: MAC addresses; multiple access protocols; some Ethernet standards; Physical layer.

# Networking: The full picture

Kurose and Ross enumerated the number of steps that are taken 'behind the scenes' when a student connects their laptop to the university network *via* an **ethernet cable**.

- Using Wi-Fi would be similar.

They list **24** steps in total ...!

A summary is useful at this stage to see how what we have learnt, for all of the various layers, come together for an everyday operation.

# 1. Device discovery

After connecting, the laptop needs an **IP address**:

- Laptop's OS places a **DHCP request**<sup>1</sup> into a **UDP segment**, itself placed into an **IP datagram**, which is then **broadcast** across the ethernet within a link-layer **frame**.
- Ethernet switch sends this message to all outgoing ports, including the **router** which extracts the request.
- Router returns an **DHCP message** (inside a UDP segment, inside an IP datagram, inside a link-layer frame).
- Laptop extracts DHCP message, its IP address, and the address of the DNS server it will use.

---

<sup>1</sup>Recall DHCP = Dynamic Host Configuration Protocol; cf. Lecture 17.

## 2. Getting the router's MAC address

The student now types `www.google.com` into their browser, which needs to be converted to an IP address using the DNS:

- Browser creates a **DNS query** and puts it into a **UDP segment** with a destination **port** of 53.
- Sent to the configured DNS server (as always, after placing in an IP datagram inside a link-layer frame).
- Does not yet know the MAC address of the router, so creates an **ARP query**<sup>1</sup> and broadcasts within a link-layer frame.
- Router sends an **ARP reply** (in a link-layer frame) to the laptop's OS.

---

<sup>1</sup>Recall ARP = Address Resolution Protocol; cf. Lecture 19.

### 3. Getting the IP address

The laptop can now send its DNS query *via* the router.

- Router extracts query (from the frame, datagram, segment) and uses its **forwarding table** to see where to send it.
- Places inside a segment, datagram, frame, and sends.
- Travels to the DNS server *via* various routers (RIP; OSPF; BGP; *cf.* Lecture 18).
- DNS server checks its **cache** and, presuming it is found, sends a **DNS reply** to the laptop (segment/datagram/frame).
- The Laptop's OS extracts the message and the IP address for `www.google.com`.



## 4. Downloading the web page

- Web browser creates a **TCP socket**, which first performs a **handshake with the web server**.
- Each stage in this handshake uses the port number for the browser, and port number 80 for the server.
- Browser creates an HTTP GET request which is forwarded *via* one or more router(s) to the server.
- The server returns an HTTP response message to the laptop.
- The laptop extracts the information and displays it.

Of course, all messages sent between the browser and the server are placed in a TCP segment, itself placed in an IP datagram, itself placed in a link-layer frame . . .

# Finally...

Good luck in the exam